

Open Information Extraction

Dominik Both, Tonio Weidler

Proseminar *Text Mining*
Andrea Zielinski

Institut für Computerlinguistik, Universität Heidelberg, 15.07.2016

ooo
oooooooo
oooooooo

oooo
ooooooooo
oooooooo
o

o
o

o
o

Strukturierung

- 1 Introduction to Information Extraction
- 2 OIE - Principles
- 3 Example: LODifier
- 4 OIE Systems in Context
- 5 Conclusion

ooo
oooooooo
oooooooo

oooo
ooooooooo
oooooooo
o

o
o

o
o

Introduction to Information Extraction

ooo
oooooooo
oooooooo

oooo
ooooooooo
oooooooo
o

o
o

o
o

What is Information Extraction?

Goal of Information Extraction is automatically extracting information from unseen text Information: entities, relations, events...

To make the dough for a good pizza, we start with putting 1kg of flour into the mixing bowl. (1kg of flour, put into, mixing bowl)

ooo
oooooooo
oooooooo

oooo
ooooooooo
oooooooo
o

o
o

o
o

Problems of Information Extraction

- Named Entity Recognition
- Relationship Extraction
- Coreference Resolution
- Comment Extraction
- many more

OIE - Principles

●○○
○○○○○○○
○○○○○○○

○○○○
○○○○○○○○○
○○○○○○○
○

○
○

○
○

OIE - Principles

Open Information Extraction

Open Information Extraction

IE: Extractor for each target relation
Open: No pre-specified extractors
Unsupervised learning of relation phrases
Extraction of information on every given domain



Problems of Open Information Extraction

- Incoherent extractions:
- This guide contains dead links and omits sites -> contains omits
- Uninformative extractions:
- Faust made a deal with the devil -> (Faust, made, a deal)

○○○
●○○○○○
○○○○○○○

○○○○
○○○○○○○○
○○○○○○○
○

○
○

○
○

OIE - Principles

Methods

Text Runner and WOE

- 1. Label: Automatic sentence labeling by heuristics
- 2. Learn: A relation phrase extractor is learned
- 3. Extract: Identifying NP pairs and searching relations words between

Problems

- Large number of labeled training examples required
- Alternative heuristic labeling leads to huge noise and stacked uncertainty
- Ignores both holistic and lexical aspects

○○○
 ○○○●○○○
 ○○○○○○○○

○○○○
 ○○○○○○○○
 ○○○○○○
 ○

○
 ○

○
 ○

Syntactic constraint

- Limits relations to those matching a certain POS Tag pattern:
- $V \mid VP \mid VW^*P$
- Always chooses longest possible match
- Merge adjacent matches together

$V \mid VP \mid VW^*P$
 V = verb particle? adv?
 W = (noun | adj | adv | pron | det)
 P = (prep | particle | inf. marker)

Lexical constraint

- Only assume relations that appear in the corpus for a certain amount
- The Obama administration is **offering only modest greenhouse gas reduction targets** at the conference


```

ooo
oooooooo●
ooooooooo

```

```

oooo
ooooooooo
ooooooooo
ooooooo
o

```

```

o
o

```

```

o
o

```

ReVerb Extraction Algorithm

- Relation Extraction: Find the longest possible string of words that match the relation constraints, merge adjacents
- Argument Extraction: Find the nearest NP left and right to the relation that is not a relativ pronoun, WHO-adverb or existential-there.
- How is the lexical constraint being checked? By creating a list of relational phrases by applying this algorithm on a 500 million Web sentences.

ooo
 ooooooooo
 ooooooooo

oooo
 oooooooooo
 ooooooooo
 o

o
 o

o
 o

ReVerb Confidence Function

- The Algorithm has a high recall, but low precision
- Now the extracted relation is weighted by a confidence function:

Weight	Feature
1.16	(x, r, y) covers all words in s
0.50	The last preposition in r is <i>for</i>
0.49	The last preposition in r is <i>on</i>
0.46	The last preposition in r is <i>of</i>
0.43	$len(s) \leq 10$ words
0.43	There is a WH-word to the left of r
0.42	r matches VW*P from Figure 1
0.39	The last preposition in r is <i>to</i>
0.25	The last preposition in r is <i>in</i>
0.23	$10 \text{ words} < len(s) \leq 20 \text{ words}$
0.21	s begins with x
0.16	y is a proper noun
0.01	x is a proper noun
-0.30	There is an NP to the left of x in s
-0.43	$20 \text{ words} < len(s)$
-0.61	r matches V from Figure 1
-0.65	There is a preposition to the left of x in s
-0.81	There is an NP to the right of y in s
-0.93	Coord. conjunction to the left of r in s

ooo
oooooooo
●oooooooo

oooo
oooooooooo
ooooooooo
o

o
o

o
o

OIE - Principles

Data Representation

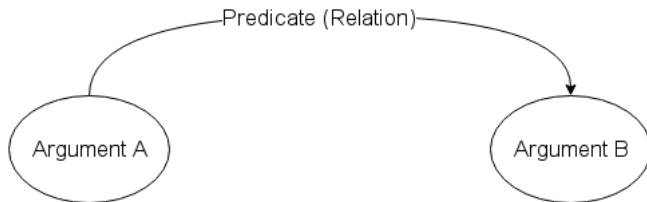
ooo
oooooooo
o●oooooooo

oooo
oooooooooo
ooooooooo
o

o
o

o
o

Standard Patterns



Argument A is in a directed **relation** to **Argument B**.

ooo
oooooooo
oo●oooo

oooo
ooooooooo
oooooooo
o

o
o

o
o

Unnormalized Annotation

(argument_a, predicate_x, argument_b)
(argument_a, predicate_y, argument_c)
(argument_a, predicate_y, argument_d)

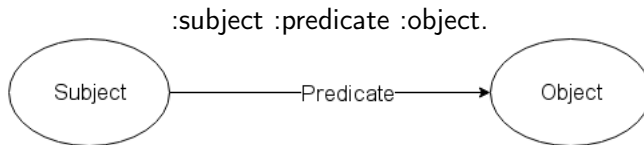
Problems

- redundant
- unnormalized
- can only produce binary predicates

RDF and Linked Data

Resource Description Framework

Models propositions by constructing *triples* including **Subjects**, **Objects** and **Predicates**
Generates a directed graph



ooo
oooooooo
oooo●ooo

oooo
ooooooooo
oooooooo
o

o
o

o
o

RDF Concepts and Notation

- **URIs**
identifies ressources (S, R, O) distinctivly and references further informations (triples)
- **Conclusions**
allows to draw conclusions using rules
- **Turtle**
allows syntax abbreviations
- **Queries**
can be searched by querying (eg SPARQL)

ooo
oooooooo
oooooooo●oo

oooo
ooooooooo
ooooooooo
o

o
o

o
o

Vocabularies & Ontologies

Several vocabularies provide useful relations and functionality, eg.:

- RDF (rdf:type, ...)
- RDFS (rdfs:subClassOf, rdfs:domain, rdfs:range, ...)
- OWL (owl:sameAs, owl:SymmetricProperty, ...)
- FOAF

Ontologies are huge RDF Graphs containing many triples, eg.:

- DBpedia
- Wikidata
- WordNet

RDF Syntax

```
dbr:Barack_Obama a foaf:person, :President;  
                  dbo:spouse dbr:Michelle_Obama.  
dbr:Bernie_Sanders dbo:birthPlace dbr:New_York,  
                                     dbr:Brooklyn;  
dbr:Brooklyn dbo:isPartOf dbr:New_York
```


ooo
oooooooo
oooooooo●


oooo
ooooooooo
ooooooooo
ooooo
o

o
o

o
o

Data Representation

... as Graph



ooo

oooooooo
oooooooo

oooo

oooooooooo
oooooooo
oooooo
o

o

o

o

o

o

Example: LODifier

ooo
oooooooo
oooooooo

oooo
ooooooooo
oooooooo
o

o
o

o
o

LODifier: Generating Linked Data from Unstructured Text (Augenstein et al., 2012)

Generate an RDF Graph from unstructured Text

Past Approaches: Use Patterns to trade recall for precision

LODifier: Process the entire text

○○○

○○○○○○○

○○○○○○○○

●○○○

○○○○○○○○○

○○○○○○○

○

○

○

○

○

○

Example: LODifier Architecture

○○○
○○○○○○○
○○○○○○○

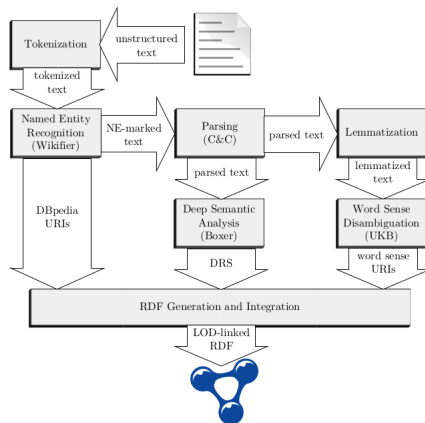
○●○○○
○○○○○○○
○○○○○○○
○

○
○

○
○

Architecture

Architecture





Approach

- 1 **Parse** the input text (POS, Treetagging, NER)
- 2 Apply **Deep Semantic Analysis** to get relations
- 3 Enrich NEs and words with **URIs** (DBpedia and WordNet)
- 4 Forge an **RDF Graph** of this information

ooo
oooooooo
oooooooo

ooo●
oooooooo
oooooooo
o

o
o

o
o

How does it happen?

Lets go through the process step-by-step!

Example Text:

The New York Times reported that John McCarthy died. He invented the programming language LISP.

example taken from Augenstein et al., 2012

○○○
○○○○○○○
○○○○○○○○

○○○○
●○○○○○○○
○○○○○○○
○

○
○

○
○

Example: LODifier

Preprocessing

Named Entity Recognition - Wikifier

Wikifier

Recognizes NE and replaces them with the Wikipedia Page Link
Disambiguates by comparing links between pages.

Example Text Output:

[The New York Times] reported that [John McCarthy (computer scientist)|John McCarthy] died. He invented the [Programming language|programming language] [Lisp (programming language)|LISP].



Parsing Syntax - C&C

C&C Parser

Syntactical Parser that tags POS and builds Parse Trees (CCG).

```

ooo
ooooooo
ooooooo

```

```

oooo
oooo●oooo
ooooooo
o

```

```

o
o

```

```

o
o

```

Preprocessing

Parsing - Output

```

ccg(1, rp(s:dcl,
  ba(s:dcl,
    lx(np, n,
      t(n, 'The_New_York_Times', 'The_New_York_Times', 'NNS', 'I-NP', '0')),
    fa(s:dcl\np,
      t((s:dcl\np)/s:em, 'reported', 'report', 'VBD', 'I-VP', '0'),
      fa(s:em,
        t(s:em/s:dcl, 'that', 'that', 'IN', 'I-SBAR', '0'),
        ba(s:dcl,
          lx(np, n,
            t(n, 'John_McCarthy', 'John_McCarthy', 'NNP', 'I-NP', 'I-PER')),
            t(s:dcl\np, 'died', 'die', 'VBD', 'I-VP', '0')))),
      t(period, '.', '.', '.', '0', '0'))).
ccg(2, rp(s:dcl,
  ba(s:dcl,
    t(np, 'He', 'he', 'PRP', 'I-NP', '0'),
    fa(s:dcl\np,
      t((s:dcl\np)/np, 'invented', 'invent', 'VBD', 'I-VP', '0'),
      fa(np:nb,
        t(np:nb/n, 'the', 'the', 'DT', 'I-NP', '0'),
        fa(n,
          t(n/n, 'programming_language', 'programming_language', 'NN', 'I-NP', '0'),
          t(n, 'LISP', 'LISP', 'NNP', 'I-NP', '0')))),
      t(period, '.', '.', '.', '0', '0'))).

```

○○○
○○○○○○○
○○○○○○○○○

○○○○
○○○○●○○○
○○○○○○○
○

○
○

○
○

Find Relations - Boxer

Boxer

Creates DRSs from C&C Output



Find Relations - Boxer

Boxer

Creates DRSs from C&C Output

Discours Representation Structure (DRS)

Represents the discourse via *relations* between *entities*

Allows referencing over the entire discourse



Find Relations - Boxer

Boxer

Creates DRSs from C&C Output

Discours Representation Structure (DRS)

Represents the discourse via *relations* between *entities*
Allows referencing over the entire discourse

Boxers DRS Relations (Conditions):

- **Unary Relations (Classes):** eg. *topic*, *person*, *event*, *male*, ...
+ all verbs
- **Binary Relations:** agent, patient, ... (semantic roles)

○○○
 ○○○○○○
 ○○○○○○


○○○
 ○○○○○○
 ○○○○○○
 ○

○
 ○

○
 ○

Preprocessing

Boxer Output





Assign WordNet URIs

RDF WordNet

WN: Lexicography containing senses linked by semantic relations

RDF WN: LD Representation of WN providing URIs for words

Steps:

- 1 Lemmatization
- 2 WSD (UKB)
- 3 Assign RDF WN URIs to word senses

○○○
○○○○○○○
○○○○○○○

○○○○
○○○○○○○●
○○○○○○○
○

○
○

○
○

Preprocessing Result

We now have ...

- URIs for all NEs
- URIs for all (disambiguated) words
- Relations between entities (those URIs)

○○○
○○○○○○○
○○○○○○○

○○○○
○○○○○○○○
●○○○○○
○

○
○

○
○

Example: LODifier

RDF Construction

○○○
○○○○○○○
○○○○○○○

○○○○
○○○○○○○○
○●○○○○○
○

○
○

○
○

What now?

Let's now construct the RDF Graph from this information!



Namespaces/Vocabularies

LODifier introduces several namespaces:

- **drsclass:** contains Boxer classes (event, person, ...) and :named relation
- **drsrel:** contains Boxer relations (agent, patient, ...)
- **ne:** contains the named entity URIs
- **reify:**



Namespaces/Vocabularies

And uses standard namespaces:

- **rdf:** mainly for `rdf:type`
- **owl:** for `owl:sameAs`

ooo
oooooooo
oooooooo

oooo
oooooooo
oo●oooo
o

o
o

o
o

Namespaces/Vocabularies

Finally the two ontologies:

- **wn30**: contains all WordNet URIs
- **dbpedia**: contains the dbpedia URIs
- **class**: contains classes not in wn30 nor in dbpedia



RDF Construction Strategy

- 1 Create a blanknode `_:x` for each discourse referent (`x0`, `x1`, ...)



RDF Construction Strategy

- 1 Create a blanknode `_:x` for each discourse referent (x_0, x_1, \dots)
- 2 if NE, then create
`_:x drsclass:named ne:URI`



RDF Construction Strategy

- 1 Create a blanknode `_:x` for each discourse referent (`x0`, `x1`, ...)
- 2 if NE, then create
`_:x drsclass:named ne:URI`
- 3 if NE and DBpedia URI exists create
`_:x owl:sameAs dbpedia:URI`



RDF Construction Strategy

- 1 Create a blanknode `_:x` for each discourse referent (`x0`, `x1`, ...)
- 2 if NE, then create
`_:x drsclass:named ne:URI`
- 3 if NE and DBpedia URI exists create
`_:x owl:sameAs dbpedia:URI`
- 4 via `rdf:type` assign closed classes (event, person, ...)
`_:x rdf:type drsclass:CLOSEDCLASS`



RDF Construction Strategy II

- 1 via `rdf:type` assign open classes (die, programming_language, ...)
`_:x rdf:type wn30:OPENCLASS, class:OPENCLASS`



RDF Construction Strategy II

- 1 via `rdf:type` assign open classes (die, programming_language, ...)
`_:x rdf:type wn30:OPENCLASS, class:OPENCLASS`
- 2 create triples from binary relations (agent, theme, ...)
`_:x drsrel:RELATION _:y`



RDF Construction Strategy II

- 1 via `rdf:type` assign open classes (die, programming_language, ...)
`_:x rdf:type wn30:OPENCLASS, class:OPENCLASS`
- 2 create triples from binary relations (agent, theme, ...)
`_:x drsrel:RELATION _:y`
- 3 REIFY

```

ooo
oooooooo
oooooooo

```

```

oooo
ooooooooo
ooooooooo
ooooooooo
o

```

```

o
o

```

```

o
o

```

RDF Construction

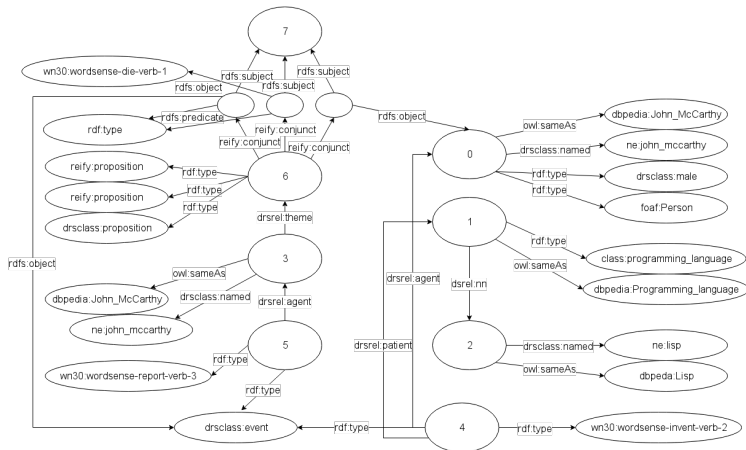
RDF Construction: Output

```

_:var0x0 drsclass:named ne:john_mccarthy ;
  rdf:type drsclass:male , foaf:Person ;
  owl:sameAs dbpedia:John_McCarthy_(computer_scientist) .
_:var0x1 rdf:type class:programming_language ;
  owl:sameAs dbpedia:Programming_language .
_:var0x2 drsrel:nn _:var0x1 .
_:var0x2 drsclass:named ne:lisp ;
  owl:sameAs dbpedia:Lisp_(programming_language) .
_:var0x3 drsclass:named ne:the_new_york_times ;
  owl:sameAs dbpedia:The_New_York_Times .
_:var0x4 rdf:type drsclass:event , wn30:wordsense-invent-verb-2 .
  drsrel:agent _:var0x0 ; drsrel:patient _:var0x2 .
_:var0x5 rdf:type drsclass:event , wn30:wordsense-report-verb-3 ;
  drsrel:agent _:var0x3 ; drsrel:theme _:var0x6 .
_:var0x6 rdf:type drsclass:proposition , reify:proposition , reify:conjunction ;
  reify:conjunct [ rdf:subject _:var0x7 ;
    rdf:predicate rdf:type ;
    rdf:object drsclass:event . ]
  reify:conjunct [ rdf:subject _:var0x7 ;
    rdf:predicate rdf:type ;
    rdf:object wn30:wordsense-die-verb-1 . ]
  reify:conjunct [ rdf:subject _:var0x7 ;
    rdf:predicate drsrel:agent ;
    rdf:object _:var0x0 . ]

```

RDF Construction: Output as Graph



ooo

oooooooo
oooooooo

oooo

oooooooooo
oooooooo

o

o

o

o

Example: LODifier

Conclusions

ooo
ooooooo
ooooooo

oooo
ooooooooo
ooooooo
o

o
o

o
o

OIE Systems in Context

ooo
oooooooo
oooooooo

oooo
ooooooooo
oooooooo
o

●
○

○
○

OIE Systems in Context

Comparison

ooo
ooooooo
ooooooo

oooo
ooooooooo
ooooooo
o

o
●

o
o

OIE Systems in Context

Evaluating the Approaches

ooo
oooooooo
oooooooo

oooo
ooooooooo
oooooooo
o

o
o

o
o

Conclusion

ooo
oooooooo
oooooooo

oooo
ooooooooo
oooooooo
o

o
o

●
o

Conclusion

Problems and Obstacles

ooo
oooooooo
oooooooo

oooo
ooooooooo
oooooooo
o

o
o

o
●

Conclusion Future Opportunities